

BankTweak: Adversarial Attack against Multi-Object Trackers by Manipulating Feature Banks

Woojin Shin¹, Donghwa Kang², Daejin Choi³, Brent Kang², Jinkyu Lee⁴, Hyeongboo Baek

¹University of Seoul, ²Korea Advanced Institute of Science and Technology(KAIST), ³Incheon National University, ⁴Sungkyunkwan University

2025.09.04
HyorinJung

Index

- **Introduction**
- **Background**
- **Analyze & Delineate**
 - **Attack formulation**
 - **How BankTweak Works: A Two-Step Process**
 - Step 1: The Groundwork (Frames t+1 to t+3)
 - Step 2: The ID Switch (Frames t+4 & t+5)
 - Solving perturbations : Using Similarity & Dissimilarity Loss
 - **Experiments**
 - Setting
 - Result
 - Ablation Study
- **Limitation & Future Direction**
- **Conclusion**

Introduction

Detection-Phase Attacks

Initial approaches aim to degrade detection quality by creating false negatives or false alarms

- Lacks Efficiency : Accuracy is only reduced in the specific frames under attack
- No Lasting Impact : The tracker recovers as soon as the attack stops

Disrupting the association phase

Recent methods manipulate object positions to cause Identity (ID) switches during the association phase.

- Lasting Impact
- Lacks Robustness : Can be easily countered by adjusting distance-related parameters
- Predictable: Relies on physically shifting the object's perceived location.

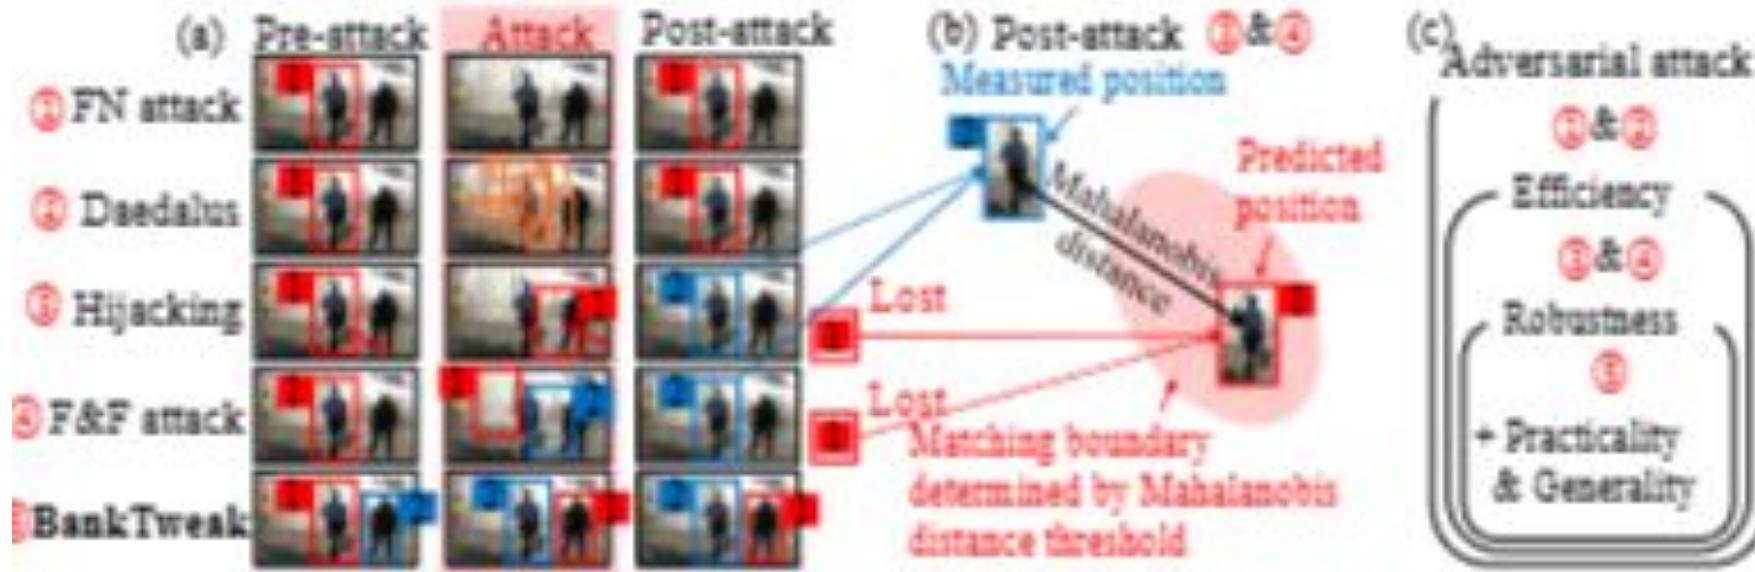
Introduction

Alternative: “Efficiency & Robustness” Solution - **BankTweak**

- It targets the **feature extractor** in the association phase
- A strategic attack that induces persistent ID switches by manipulating features in the feature bank **without altering object positions**

👉 BankTweak overcomes the limitations of previous methods by targeting the feature bank, the core of the association phase

Introduction



👉 Efficiency & Robustness & Practicality & Generality

Background

Modern Multi-Object Tracking (MOT) works in two main stages to identify objects and follow them across video frames

- **Detention**
 - A detector identifies and draws bounding boxes around all objects of interest within a single frame
- **Association**
 - Links new detections to existing tracks using appearance features and motion cues
 - Sequential matching:
 - Feature-based matching pairs objects with highest feature similarity using the feature bank
 - Motion-based IoU matching handles unmatched objects

Background

Feature bank

- A storage of appearance features for all tracked objects, used to compare and match new detections with existing tracks across frames
- Maintains historical object representations, enabling the tracker to recognize objects even after short occlusions or frame gaps



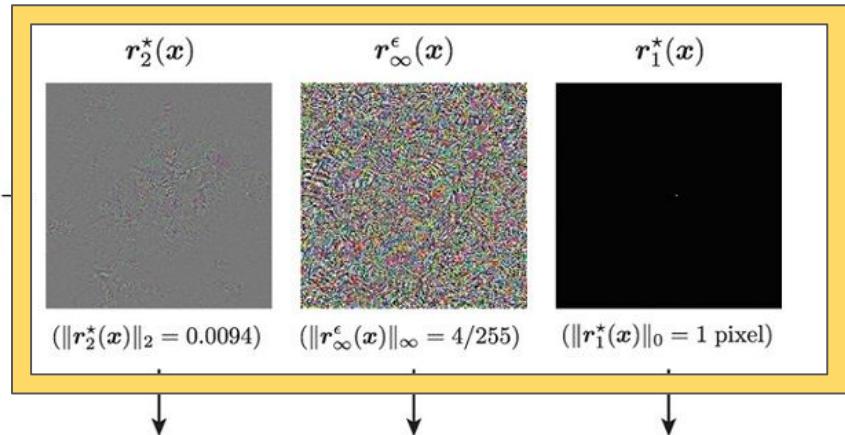
Background

Adversarial Perturbation

$$\delta = \arg \min_{\delta, \|\delta\|_\infty < \epsilon} \mathcal{L}(E(D(I + \delta|\theta_D), I + \delta|\theta_E), \mathbb{F}),$$



Persian cat



Broccoli



Sulphur butterfly



Broccoli

👉 Find perturbation δ that minimizes loss while keeping change imperceptible



Attack formulation

PGD(Projected Gradient Descent)

$$\begin{aligned}\delta^{r+1} &= \text{clip}_{[-\epsilon, \epsilon] \cap [-I, 1-I]} \\ &(\delta^r + \alpha \text{sgn}(\nabla_{\delta} \mathcal{L}(E(D(I + \delta|\theta_D), I + \delta|\theta_E), \mathbf{F})))\end{aligned}$$

- α (alpha): Step size (per-pixel change)
- ϵ (epsilon): Max allowed perturbation (ℓ_∞ -norm bound)
- ∇ (nabla): Gradient of loss LLL w.r.t perturbation
- $\text{sgn}(\cdot)$: Uses only the sign (direction) of the gradient

Attack formulation

PGD(Projected Gradient Descent) Steps

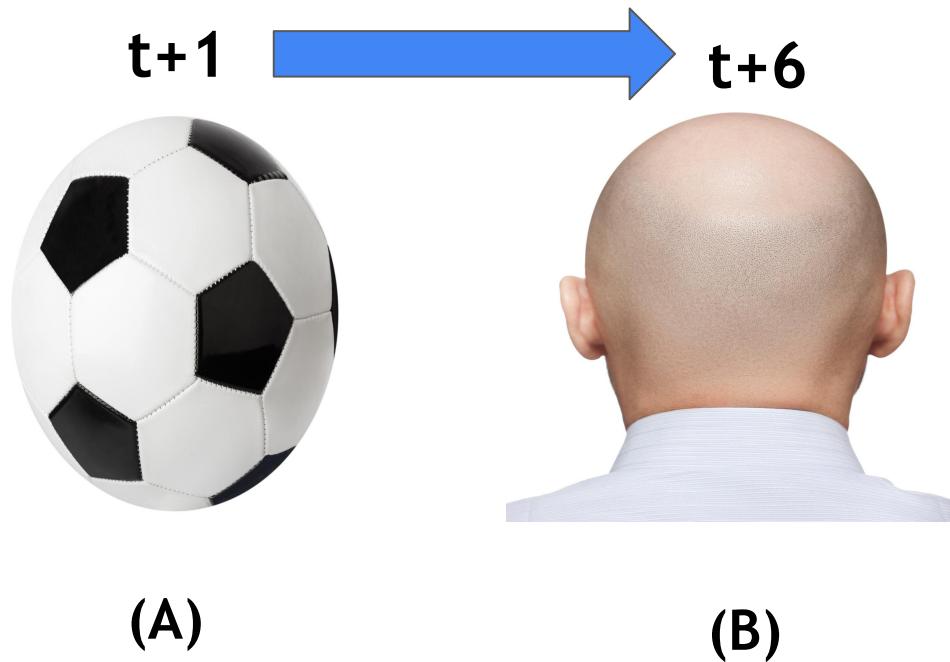
- Initialization: $\delta=0$
- Iterative Update (R steps):
 - Compute loss L on current perturbed image
 - Compute gradient ∇L
 - Update perturbation:
$$\delta_{r+1} = \delta_r + \alpha \cdot \text{sgn}(\nabla L)$$
 - Clipping:
 - Keep $|\delta| \leq \epsilon$
 - Ensure valid pixel range $[0,1]$
- Final Adversarial Image: $I' = I + \delta$

Outcome

- Imperceptible to humans
- Disrupts detector & feature extractor
- Effective with only 5 frames

How BankTweak Works: A Two-Step Process

The attack unfolds over five consecutive frames to induce a **permanent** ID switch between two objects, (A) and (B)



Step 1: The Groundwork (Frames t+1 to t+3)

The goal is to secretly inject carefully crafted features into the feature banks of target objects (A) and (B) without triggering an ID switch prematurely.

Frame t+1	Frame t+2	Frame t+3
Inject 'dummy' features ($X A$ and $Y B$) that are very dissimilar to the original features of A and B.	Inject a feature that looks like B but is generated from A ($B A$) into object A's feature bank.	Inject a feature that looks like A but is generated from B ($A B$) into object B's feature bank.

Outcome of Step 1

At the end of this phase, the feature banks for both objects are 'poisoned' with misleading information, setting the stage for the final switch.

Step 2: The ID Switch (Frames t+4 & t+5)

With the groundwork laid, BankTweak now exploits the Hungarian matching algorithm, which allocates IDs based on the lowest average feature distance.

Frame t+4: The Switch

A specific feature is crafted for object A. The algorithm now finds it 'cheaper' to match the current object A with the history of object B (which contains the poisoned B|A feature). This confusion forces the ID switch.

Frame t+5: Solidification

A similar process is repeated for object B to solidify the swap and prevent it from reverting in the next frame.

Result: A Permanent Switch

From frame t+6 onwards, the objects' original features continue to match with the now-swapped ID histories. The ID switch is permanent, even though no more attacks are being performed

Why Not Just Swap Features Directly?

The Tracker is Too Smart

- Direct attack of injecting B's feature into A's track will fail.
- The feature bank for A still contains many instances of A's true feature. The matching algorithm will correctly prioritize the strong history of true 'A' features, causing any temporary ID switch to revert immediately.

👉 BankTweak's meticulous two-step process is necessary to overcome this resilience.

Solving perturbations : Using Similarity & Dissimilarity Loss

Loss Functions

Similarity Loss

$$\mathcal{L}^s(\mathbb{F}^*, \mathbb{F}) = \sum C(F_i^*, F_i),$$

Make current feature F_i^*
similar to target F_i
(minimize cosine distance)

Dissimilarity Loss

$$\mathcal{L}^d(\mathbb{F}^*, \mathbb{F}) = - \sum_{F_i^* \in \mathbb{F}^*, F_i \in \mathbb{F}} C(F_i^*, F_i).$$

Make current feature F_i^*
different from target F_i
(maximize cosine distance)

Solving perturbations : Using Similarity & Dissimilarity Loss

Stepwise Attack



Step 1: Groundwork

Step 2: ID Switch

- Alternate L_s and L_d across frames
- Inject manipulated features into the feature bank
- Exploit Hungarian matching vulnerability
→ achieve **permanent** ID switch

Experiments : Setting

Efficiency: IDF1, HOTA, ID Switches (IDsw)

- **HOTA:** $\sqrt{(\text{DetA} \times \text{AssA})}$
 - DetA = accurate detection ratio
 - AssA = correctly tracked object ratio
- Accuracy excludes attack frames

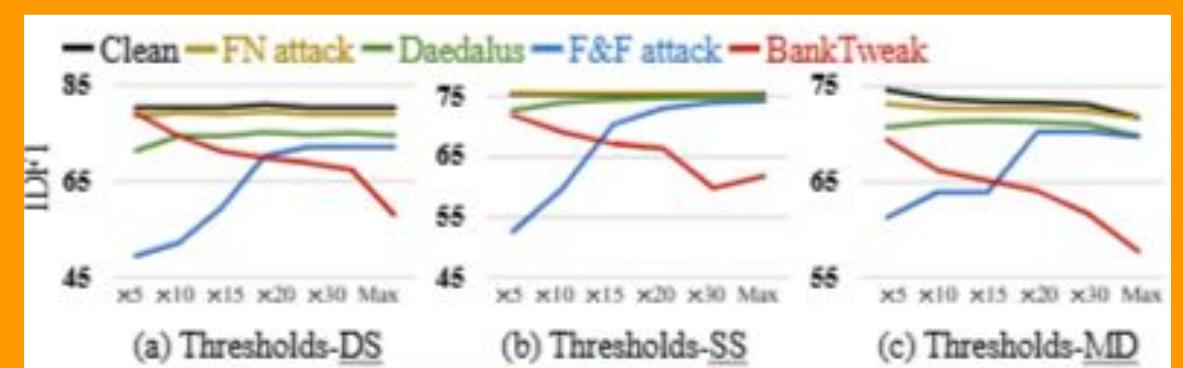
Robustness: Measured by varying Mahalanobis distance threshold

Practicality: ρ_{Det} & ρ_{ID}

- $\rho_{\text{Det}} = \text{avg. increase in detections per attack frame} / \text{GT}$
- $\rho_{\text{ID}} = \text{increase in ID counts}$

Generality: No metric; BankTweak designed to be model-agnostic

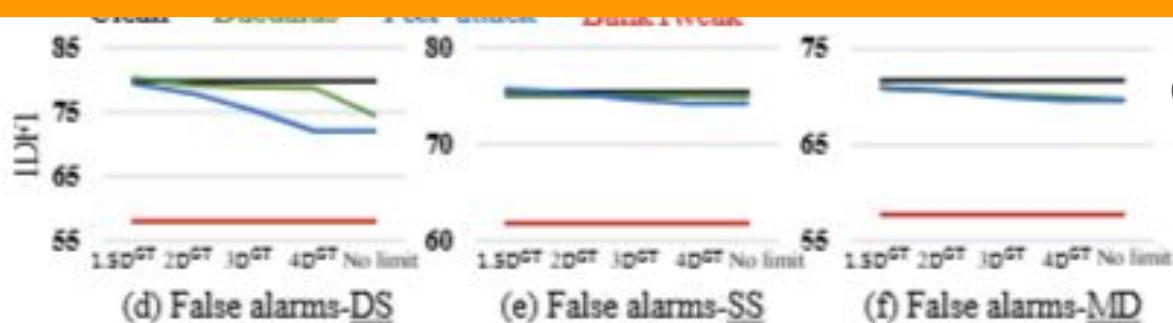
Experiments : Results



Mahalanobis Threshold

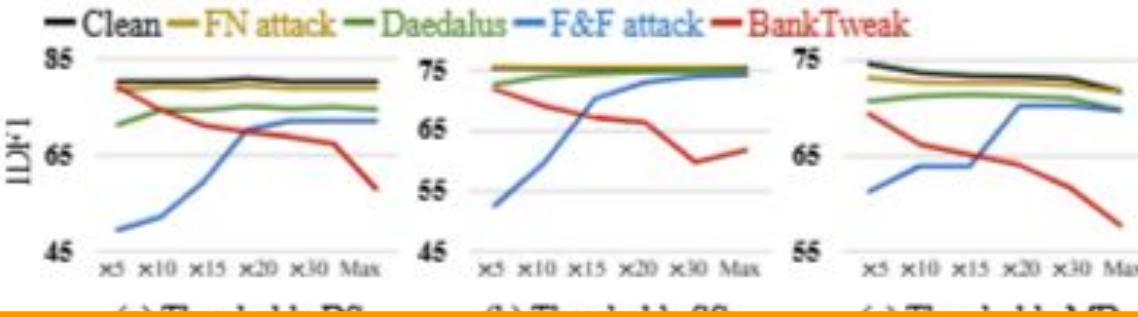
Clean & FN: IDF1 stable, unaffected by threshold

F&F: Attack weakens as threshold increases



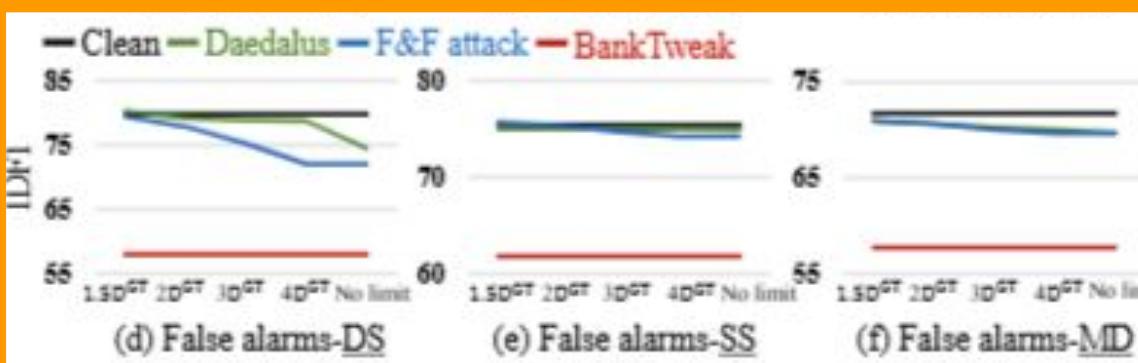
BankTweak: Low IDF1, robust to threshold changes

Experiments : Results



Allowed False Positives

Daedalus & F&F: IDF1 drops as allowed false positives increase, but remains higher than BankTweak



⚠️ **Clean & BankTweak:** IDF1 stable regardless of false positive limit

Experiments: Ablation study

Mahalanobis Threshold

- F&F, Daedalus: \uparrow threshold $\rightarrow \downarrow$ effect
- BankTweak: robust

False Positive Limit

- Clean & BankTweak: stable (no FP)
- Daedalus & F&F: more FP \rightarrow lower IDF1, still $<$ BankTweak

Step 2

- Removed: 79.78 \rightarrow 77.52 (small drop)
- Included: 79.78 \rightarrow 58.01 (large drop)

Iterations (Rs)

- More Rs \rightarrow stronger attack
- Larger ϵ \rightarrow fewer Rs needed

Limitation & Future Directions

- The current model is a white-box attack, requiring internal knowledge of the tracker's models
- BankTweak's transfer-based black-box attacks are feasible when two target objects have **similar features**, allowing ID switches without full model access
- Future work will explore transfer-based black-box attacks, which show initial feasibility.
- The attack requires at least two objects, but this is not a practical limitation in real-world tracking scenarios where scenes are rarely empty.

Conclusion

- BankTweak induces persistent ID switches in multi-object trackers.
- Robust against changes in Mahalanobis distance thresholds.
- Does not generate false positives, ensuring practical applicability.
- Step 2 and iterative updates are crucial for attack efficiency.
- Transfer-based black-box attacks are feasible with similar-feature object pairs.
- Searching similar object pairs is computationally intensive → motivates lightweight models.